

CONTENTS

1. Acknowledgment
2. Problem Statement
3. Libraries Used
4. Exploratory Data Analysis
5. Q1) Most common emotions/words used by males and females.
6. Q2) Which gender makes more typos in their tweets?
7. Classification of the dataset
8. Final Result

Acknowledgement

IT GIVES US IMENSE PLEASURE IN BRINGING OUT THIS SYNOPSIS OF THE PROJECT ENTITLED

“ML Project on Information.csv”

FIRSTLY, WE WOULD LIKE TO THANK OUR TEACHER AND MENTOR MR. AQIB AHMED WHO GAVE US HIS VALUABLE SUGGESTIONS AND IDEAS WHEN WE WERE IN NEED OF THEM. HE ENCOURAGED US TO WORK ON THIS PROJECT.

WE ARE ALSO GRATEFUL TO VERZEO FOR GIVING US THE OPPORTUNITY TO WORK WITH THEM AND PROVIDING US THE NECESSARY RESOURCES FOR THE PROJECT

WE WOULD ALSO THANK TO ALL OF THEM WHO HELPED US TO COMPLETE THIS PROJECT.

WE ARE IMMENSELY GRATEFULL TO ALL INVOLVED IN THIS PROJECT AS WITHOUT THEIR INSPITRATION AND VALUABLE SUGGESTION IT WOULD NOT HAVE BEEN POSSIBLE TO DEVELOP THE PROJECT WITHIN THE PRESCRIBED TIME.

WITH SINCERE THANKS,

“ML062B12”

Problem Statement for Project Analysis:

For a given dataset,

Information.csv

Information.csv																
A	B	C	D	E	F	G	H	I	J	K	L	M	N	O		
unit_id	_golden	_unit_state	_trusted_judgments	_last_judgment_at	gender	gender:confidence	profile_yr	profile_yr:confidence	created	description	fav_number	gender_gold	link_color	name		
815719226	FALSE	finalized		3	10/26/15 23:24	male		1	yes	1	12/5/13 1:48	I sing my own rhythm	0	08C2C2	#heez0	
815719227	FALSE	finalized		3	10/26/15 23:30	male		1	yes	1	10/1/12 13:51	I'm the author of nov	68	0084B4	DavidBumet	
815719228	FALSE	finalized		3	10/26/15 23:33	male		0.8625	yes	1	11/28/14 11:30	louiswhining and sq	7696	AB88C2	hfpreflylaugh	
815719229	FALSE	finalized		3	10/26/15 23:10	male		1	yes	1	6/11/09 22:39	Mobile guy. 49ers, S	202	0084B4	douggaftand	
815719230	FALSE	finalized		3	10/27/15 1:15	female		1	yes	1	4/16/14 13:23	Ricky Wilson The Be	37318	3B94D9	WillfordGemma	
815719231	FALSE	finalized		3	10/27/15 1:47	female		1	yes	1	3/11/10 18:14	you don't know me.	3901	F5ABBS	monnevious	
815719232	FALSE	finalized		3	10/27/15 1:57	brand		1	yes	1	4/24/08 13:03	A global marketplac	4122	298AAE	Shutterstock	
815719233	FALSE	finalized		3	10/26/15 23:48	male		1	yes	1	12/3/12 21:54	The secret of getting	80	0000FF	RobinMeke	
815719234	FALSE	finalized		3	10/27/15 1:52	female		1	yes	1	9/8/15 4:50	Pili Fan // Crazy abou	1825	9266CC	plg3llie_	
815719235	FALSE	finalized		3	10/27/15 1:49	female		1	yes	1	5/13/11 3:32	Renaissance art hist	3115	9266CC	GabrieleNaher	
815719236	FALSE	finalized		3	10/26/15 23:17	brand		0.7002	yes	1	11/16/11 17:14	Clean food that taste	516	0084B4	Blissful_Eats	
815719237	FALSE	finalized		3	10/26/15 22:33	brand		1	yes	1	2/22/15 20:06	highly extraordinary	0	0084B4	eliseolecteo	
815719238	FALSE	finalized		3	10/26/15 22:20	female		0.6509	yes	1	8/10/12 5:05	Senior'16. XI-XII-M	3371	0084B4	kaylanabrewer	
815719239	FALSE	finalized		3	10/26/15 23:29	brand		1	yes	1	5/1/12 22:14	Come join the fastest	0	2FC2EF	americacredit	
815719240	FALSE	finalized		3	10/27/15 1:29	female		0.6501	yes	1	4/6/13 15:31	Im just here for t**p	13928	0084B4	cheyfan	
815719241	FALSE	finalized		3	10/27/15 1:50	female		1	yes	1	10/3/15 21:32		0	0084B4	Ayu74th4	
815719242	FALSE	finalized		3	10/26/15 23:43	female		1	yes	1	8/27/11 9:42	JMKM*_s**	1762		58185	Toucan_Sam
815719243	FALSE	finalized		3	10/26/15 22:50	male		1	yes	1	10/18/09 11:41	Over enthusiastic F1	5	0084B4	SSmyh2010	
815719244	FALSE	finalized		3	10/27/15 1:42	male		1	yes	1	7/20/15 12:01		1	0084B4	DarkToonsGam	
815719245	FALSE	finalized		3	10/26/15 22:19	unknown		0.3527	yes	1	1/30/15 9:52		160	0084B4	GoutayrLynn	
815719246	FALSE	finalized		3	10/27/15 1:21	female		1	yes	1	2/28/13 3:04	Artisan specializing i	18751	3B94D9	jpeter37	
815719247	FALSE	finalized		3	10/27/15 0:09	female		1	yes	1	10/14/11 17:53	He bled and died to	5454	D41EBF	SarahMaddy	
815719248	FALSE	finalized		3	10/26/15 22:47	female		1	yes	1	2/15/15 6:41	union j xxxx	133	0084B4	bethlaia	
815719249	FALSE	finalized		3	10/27/15 1:34	male		1	yes	1	1/11/13 1:18	You had me from the	2871	3B94D9	SticilyAd	
815719250	FALSE	finalized		3	10/26/15 23:01	male		1	yes	1	4/21/11 12:21	BSc economics grad	228	0084B4	Jkurlett	
815719251	FALSE	finalized		3	10/26/15 22:54	female		1	yes	1	8/14/13 12:54	Wife to my Coach. M	1051	0084B4	karenbullard	
815719252	FALSE	finalized		3	10/26/15 22:24	brand		0.6576	yes	1	8/8/15 16:15	If you have any ques	0	0084B4	1oneonlyone1	
815719253	FALSE	finalized		3	10/27/15 1:22	brand		0.6667	yes	1	4/26/15 8:15	14. Canadian. Spac	5601	D02E44	the_official_jos	

Find out the best algorithm as per accuracy.

Analysis:

After having a brief look, we found that the dataset (information.csv) gives very detailed and vast information about the tweets posted by both males and females. We also found that there was still some scope for cleansing of the data so as to classify the data more efficiently and accurately using various classification algorithms.

Thus, We went ahead and performed EDA in order to clean the dataset, in the first step of this project work.

Libraries Imported and Used for the project:

1. Numpy

2. Pandas

3. Sklearn

- Preprocessing
- Model_Selection
- Metrics
- DictVectorizer
- TfidfVectorizer
- LogisticRegression
- KNeighborsClassifier
- SVC (sklearn.svm)
- RandomForestClassifier

4. Nltk

- Corpus

5. Re

6. Matplotlib.Pyplot

7. TextBlob

MAJOR PROJECT

Problem Statement : For a given dataset (problem) which is the best classification algorithm (as per accuracy)

```
In [1]: import numpy as np
import pandas as pd
from sklearn import preprocessing
import nltk
nltk.download("stopwords")
from nltk.corpus import stopwords
import re

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\KIIT\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
In [2]: import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
In [3]: #Reading the dataset
df = pd.read_csv('Information.csv',encoding = "latin-1")
df
```

Exploratory Data Analysis (Cleaning of the Data):

The dataset provided contains: 20050 rows \times 26 columns, it needed cleaning. So the next thing we did was to perform EDA on the dataset. We needed dataset with only 'male' and 'female' as 'gender' values. So we used `pd.concat()` to remove the unnecessary values in 'gender'. Now we were left with: 12894 rows \times 26 columns. Accordingly, we used appropriate functions to clean the dataset. In addition, 'text' is cleaned and stored as 'Tweets' whereas 'description' is cleaned and stored as 'Desc'. We also dropped those rows where 'location' was not available, to be left with: 8747 rows \times 28 columns and hence, performed EDA successfully.

Here we cleaned the data in the following steps:

- Kept Data where gender was either male or female.

EDA

```
In [4]: #Keeping those data where gender is male or female
df_male = df[df["gender"] == "male"]
df_female = df[df["gender"] == "female"]
df = pd.concat([df_male, df_female])
```

- Defined a function `clean()` which uses regular expressions to clean the data.

```
In [5]: #Function to clean the dataset
def clean(s):
    s = str(s)
    s = s.lower()
    s = re.sub('\s\w', ' ', s)
    s = re.sub('\w\s', ' ', s)
    s = re.sub(r'[\^w]', ' ', s)
    s = re.sub("\d+", "", s)
    s = re.sub('\s+', ' ', s)
    s = re.sub('!@#$', '', s)
    s = s.replace("co", "")
    s = s.replace("https", "")
    s = s.replace(" ", "")
    s = s.replace("[\w*", " ")
    return s
```

- Using the `clean` function cleaned and changed words like "text" and "description" into "Tweets" and "Desc" respectively.

```
In [6]: df['Tweets'] = [clean(s) for s in df['text']]
df['Desc'] = [clean(s) for s in df['description']]

stop = set(stopwords.words('english'))
df['Tweets'] = df['Tweets'].str.lower().str.split()
df['Tweets'] = df['Tweets'].apply(lambda x : [item for item in x if item not in stop])

for i in range(df.shape[1]):
    df[df.columns[i]] = [clean(s) for s in df[df.columns[i]]]

# 'text' is cleaned and stored as 'Tweets'
# 'description' is cleaned and stored as 'Desc'
```

- Dropped data where Tweet Location was not available.

```
In [7]: #Dropped those where location was not available
df = df[df["tweet_location"] != "nan"]
df
```

Q1) What are the most common emotions/words used by Males and Females according to the given dataset?

Using Pandas' Series Function and plotting the graphs of tweet word counts of both Males and Females (using Matplotlib).

We found out that:

The most common emotions/words used by Males : ú , get

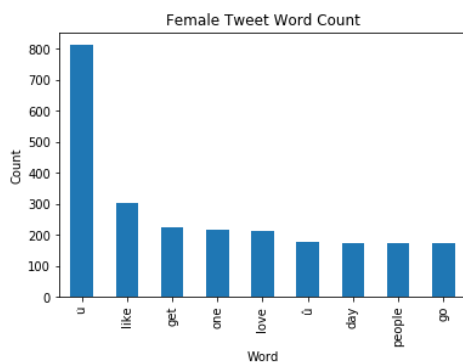
The most common emotions/words used by Females : ú , like

Q1. What are the most common emotions/words used by Males and Females?

```
In [8]: male = df[df['gender'] == 'male']
female = df[df['gender'] == 'female']
male_words = pd.Series(' '.join(male['Tweets'].astype(str)).lower().split(" ")).value_counts()[:10]
female_words = pd.Series(' '.join(female['Tweets'].astype(str)).lower().split(" ")).value_counts()[:10]
male_words = male_words.iloc[1:]
female_words = female_words.iloc[1:]
```

```
In [9]: plt.title("Female Tweet Word Count")
plt.xlabel("Word")
plt.ylabel("Count")
female_words.plot.bar()
```


Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x22a186de548>



(Female Tweet Word Count)

```
In [10]: plt.title("male Tweet Word Count")
plt.xlabel("Word")
plt.ylabel("Count")
male_words.plot.bar()
```

Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x22a175fbf88>



Answer:
The most common emotions/words used by Males : ú , get
The most common emotions/words used by Females : ú , like

(Male Tweet Word Count)

Q2) Which gender makes more typos in their tweets?

Again, Using Pandas' Series function and Python's TextBlob library over the dataset, we calculated which gender makes more Typos(Mistakes) while writing their tweets.

We found out that it was MALES who made more number of typos in their tweets as compared to FEMALES.

Q2. Which gender makes more typos in their tweets?

```
In [11]: male_words = pd.Series(' '.join(male['Tweets'].astype(str)).lower().split(" "))
female_words = pd.Series(' '.join(female['Tweets'].astype(str)).lower().split(" "))
```

```
In [12]: from textblob import TextBlob
```

```
In [13]: d=male_words.values
a=0
for word in d:
    b = TextBlob(word)
    c=str(b.correct())
    if word==c:
        pass
    else:
        a=a+1

print("Total no. of male words = ",male_words.count())
print("Total no. typos by male = ",a)
```

```
Total no. of male words = 50967
Total no. typos by male = 7590
```

```
In [14]: d=female_words.values
e=0
for word in d:
    b = TextBlob(word)
    c=str(b.correct())
    if word==c:
        pass
    else:
        e=e+1

print("Total no. of female words = ",female_words.count())
print("Total no. typos by female = ",e)
```

```
Total no. of female words = 47482
Total no. typos by female = 6599
```

```
In [15]: if (a/(male_words.count())) > (e/(female_words.count())):
    print("Males make more typos.")
else:
    print("Females make more typos.")
```

```
Males make more typos.
```

Answer:
Males make more typos.

Final Step: Classification of entire Dataset (Information.csv)

In order to get started with the classification of the dataset, we decided to take “Gender” as a Dependent Variable and “Description” as an Independent Variable, while also performing the Vectorization and Label Encoding on the data using Tfidfvectorizer() giving us the signal to finally apply classification algorithms on the data set.

Here, we are taking "gender" as our dependent variable and "description" as our independent variable.

```
In [16]: > import collections
          > from sklearn.feature_extraction import DictVectorizer
          > from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [17]: > #Performing Vectorization and Label Encoding
          > Vectorize = TfidfVectorizer(stop_words='english')
          > x = Vectorize.fit_transform(df['Desc'])
          > y = df.gender
          > le = preprocessing.LabelEncoder()
          > y = le.fit_transform(y.values)
```

We Split the data set into for training and testing purposes as usual.

```
In [18]: > #Splitting the dataset for training and testing
          > X_train, X_test, Y_train, Y_test = train_test_split(x, y)
```

Algorithm 1~ Logistic Regression:

Accuracy Achieved: 64.15%

1. Logistic Regression

```
In [19]: > from sklearn.linear_model import LogisticRegression
```

```
LogReg = LogisticRegression(random_state=25)
#Training the model
LogReg.fit(X_train, Y_train)
```

```
C:\Anaconda\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
in 0.22. Specify a solver to silence this warning.
FutureWarning)
```

```
Out[19]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='warn', n_jobs=None, penalty='l2',
random_state=25, solver='warn', tol=0.0001, verbose=0,
warm_start=False)
```

```
In [20]: > #Testing and predicting
          > y_pred = LogReg.predict(X_test)
```

```
In [21]: > #Checking the accuracy
          > from sklearn import metrics
          > print("Test set Accuracy: ", metrics.accuracy_score(Y_test, y_pred))
```

```
Test set Accuracy: 0.6415180612711477
```


Algorithm 2 – K Nearest Neighbor (KNN)

Accuracy Achieved : 56.33%

2. KNN

```
In [22]: from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=107)
#Training the model
knn.fit(X_train, Y_train)

Out[22]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=107, p=2,
                             weights='uniform')

In [23]: #Testing and predicting
y_pred = knn.predict(X_test)

In [24]: #Checking the accuracy
print("Test set Accuracy: ", metrics.accuracy_score(Y_test, y_pred))

Test set Accuracy: 0.563328760859625
```

Accuracy = 56.33%

Algorithm 3 – Support Vector Machine (SVM)

Accuracy Achieved : 51.53%

3. SVM

```
In [25]: from sklearn.svm import SVC

svc = SVC(kernel='rbf', random_state=25)
#Training the model
svc.fit(X_train, Y_train)

C:\Anaconda\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to
'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this war
ning.
"avoid this warning.", FutureWarning)

Out[25]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
             decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
             kernel='rbf', max_iter=-1, probability=False, random_state=25,
             shrinking=True, tol=0.001, verbose=False)

In [27]: #Testing and predicting
y_pred = svc.predict(X_test)

In [28]: from sklearn.metrics import classification_report, confusion_matrix
print (classification_report(Y_test, y_pred))

              precision    recall  f1-score   support

         0           0.00         0.00         0.00         1060
         1           0.52         1.00         0.68         1127

 accuracy          0.26
 macro avg         0.26         0.50         0.34         2187
 weighted avg      0.27         0.52         0.35         2187

C:\Anaconda\lib\site-packages\sklearn\metrics\classification.py:1437: UndefinedMetricWarning: Precision and F-score are ill-
defined and being set to 0.0 in labels with no predicted samples.
'precision', 'predicted', average, warn_for)

In [29]: #Checking the accuracy
print("Test set Accuracy: ", metrics.accuracy_score(Y_test, y_pred))

Test set Accuracy: 0.5153177869227252
```

Accuracy = 51.53%

Algorithm 4 – Random Forest

Accuracy Achieved : 60.95%

4. Random Forest

```
In [30]: from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier(random_state=25)
#Training the model
rfc.fit(X_train, Y_train)

C:\Anaconda\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)

Out[30]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10,
                                n_jobs=None, oob_score=False, random_state=25, verbose=0,
                                warm_start=False)

In [31]: #Testing and predicting
y_pred = rfc.predict(X_test)

In [32]: from sklearn.metrics import classification_report, confusion_matrix
print (classification_report(Y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.57	0.76	0.65	1060
1	0.67	0.47	0.55	1127
accuracy			0.61	2187
macro avg	0.62	0.61	0.60	2187
weighted avg	0.63	0.61	0.60	2187

```
In [33]: #Checking the accuracy
print("Test set Accuracy: ", metrics.accuracy_score(Y_test, y_pred))

Test set Accuracy:  0.6095107453132145
```

Accuracy = 60.95%

From the above results, it is found that Logistic Regression has the greatest accuracy(64.15%) in predicting the gender of the a user

Final Result:

It was found that the highest accuracy was achieved by the Logistic Regression Classification Algorithm i.e. 64.15%, thus making it the most suitable algorithm for the given dataset INFORMATION.CSV.

Thank You.